Semi-supervised protein classification using cluster kernels

Jason Weston[a],[*] *Christina Leslie*[b], *Eugene Ie*[b], *Dengyong Zhou*[c],
Andre Elisseeff[c], *William Stafford Noble*[d]

[a]*NEC Research Institute, 4 Independence Way, Princeton, NJ 08540, USA,* [b]*Center for Computational Learning Systems, Columbia University, Interchuch Center, 475 Riverside Dr., Mail Code 7717, New York, NY 10115, USA,* [c]*Max-Planck Institute for Biological Cybernetics, Spemannstraße 38, 72076 Tübingen, Germany,* [d]*Department of Genome Sciences, University of Washington, 1705 NE Pacific Street, Seattle, WA 98195, USA*

ABSTRACT
Motivation: Building an accurate protein classification system depends critically upon choosing a good representation of the input sequences of amino acids. Recent work using string kernels for protein data has achieved state-of-the-art classification performance. However, such representations are based only on labeled data—examples with known 3D structures, organized into structural classes—while in practice, unlabeled data is far more plentiful.
Results: In this work, we develop simple and scalable cluster kernel techniques for incorporating unlabeled data into the representation of protein sequences. We show that our methods greatly improve the classification performance of string kernels and outperform standard approaches for using unlabeled data, such as adding close homologs of the positive examples to the training data. We achieve equal or superior performance to previously presented cluster kernel methods while achieving far greater computational efficiency.
Availability: Supplementary data and source code are available at `www.kyb.tuebingen.mpg.de/bs/people/weston/semiprot`. The Spider matlab package is available at `www.kyb.tuebingen.mpg.de/bs/people/spider`.
Contact: jasonw@nec-labs.com

1 INTRODUCTION

A central problem in computational biology is the classification of proteins into functional and structural classes given their amino acid sequences. The major methods for homology detection can be split into three basic groups: pairwise sequence comparison algorithms (Altschul et al., 1990; Smith and Waterman, 1981), generative models for protein families (Krogh et al., 1994; Park et al., 1998), and discriminative classifiers (Jaakkola et al., 2000; Leslie et al., 2002; Liao and

Noble, 2002). Popular sequence comparison methods such as BLAST and Smith-Waterman are based on unsupervised alignment scores. Generative models such as profile hidden Markov models (HMMs) model positive examples of a protein family, but they can be trained iteratively using both positively labeled and unlabeled examples by pulling in close homologs and adding them to the positive set. A compromise between these methods is PSI-BLAST (Altschul et al., 1997), which uses BLAST to iteratively build a probabilistic profile of a query sequence and obtain a more sensitive sequence comparison score. Finally, classifiers such as support vector machines (SVMs) use both positive and negative examples and provide state-of-the-art performance when used with appropriate distance metrics (i.e., appropriate kernels) (Jaakkola et al., 2000; Leslie et al., 2002; Liao and Noble, 2002; Saigo et al., 2004). However, these classifiers still require an auxiliary method (such as PSI-BLAST) to handle unlabeled data: one generally adds predicted homologs of the positive training examples to the training set before training the classifier.

In practice, relatively little labeled data is available—approximately 30,000 proteins with known 3D structure, some belonging to families and superfamilies with only a handful of labeled members—whereas there are close to one million sequenced proteins, providing abundant unlabeled data. New semi-supervised learning techniques should be able to make better use of this unlabeled data.

Recent work in semi-supervised learning has focused on changing the representation given to a classifier by taking into account the structure described by the unlabeled data (Zhu and Ghahramani, 2002; Chapelle et al., 2002; Szummer and Jaakkola, 2001). These works can be viewed as cases of *cluster kernels*, which produce similarity metrics based on the cluster assumption: namely, two points in the same "cluster" or region of high density should have a small distance

[*]to whom correspondence should be addressed

to each other. In this work, we investigate the use of cluster kernels for protein classification by developing two simple and scalable methods for modifying a base kernel. The *neighborhood kernel* uses averaging over a neighborhood of sequences defined by a local sequence similarity measure, and the *bagged kernel* uses bagged clustering of the full sequence data set to modify the base kernel. In both the semi-supervised and transductive settings, these techniques greatly improve classification performance when used with mismatch string kernels, and the techniques achieve equal or superior results to all previously presented cluster kernel methods that we tried. Moreover, the neighborhood and bagged kernel approaches are far more computationally efficient than these competing methods.

The current work is an expanded version of a conference proceedings paper (Weston et al., 2003). We haved included new large scale experiments for the cluster kernel methods using the Swiss-Prot database as a source of unlabeled protein sequence data and comparison to the recent profile kernel method (Kuang et al., 2004).

2 REPRESENTATIONS AND KERNELS FOR PROTEIN SEQUENCES

Proteins can be represented as variable length sequences, typically several hundred characters long, from the alphabet of 20 amino acids. In order to use learning algorithms that require vector inputs, we must first find a suitable feature vector representation, mapping sequence x into a vector space by $x \mapsto \Phi(x)$. Kernel methods such as SVMs only need to compute inner products, called kernels, $K(x, y) = \langle \Phi(x), \Phi(y) \rangle$, for training and testing. We thus can accomplish the above mapping using a kernel for sequence data.

Biologically motivated sequence comparison scores, like Smith-Waterman or BLAST, provide an appealing representation of sequence data. The Smith-Waterman (SW) algorithm (Smith and Waterman, 1981) uses dynamic programming to compute the optimal local gapped alignment score between two sequences, while BLAST (Altschul et al., 1990) approximates SW by computing a heuristic alignment score. Both methods return empirically estimated E-values indicating the confidence of the score. These alignment-based scores do not define a positive definite kernel; however, one can use a feature representation based on the empirical kernel map

$$\Phi(x) = \langle d(x_1, x), \ldots, d(x_m, x) \rangle$$

where $d(x, y)$ is the pairwise score (or E-value) between x and y and x_i, $i = 1 \ldots m$, are the training sequences. Using SW E-values in this fashion gives strong classification performance (Liao and Noble, 2002). Note, however, that the method is slow, both because computing each SW score is $O(|x|^2)$ and because computing each empirically mapped kernel value is $O(m)$.

Another appealing idea is to derive the feature representation from a generative model for a protein family. In the Fisher kernel method (Jaakkola et al., 2000), one first builds a profile HMM for the positive training sequences, defining a log likelihood function $\log P(x|\theta)$ for any protein sequence x. Then the gradient vector $\nabla_\theta \log P(x|\theta)|_{\theta=\theta_0}$, where θ_0 is the maximum likelihood estimate for model parameters, defines an explicit vector of features, called Fisher scores, for x. This representation gives excellent classification results, but the Fisher scores must be computed by an $O(|x|^2)$ forward-backward algorithm, making the kernel tractable but slow.

It is possible to construct useful kernels directly without explicitly depending on generative models by using string kernels. For example, the mismatch kernel (Leslie et al., 2002) is defined by a histogram-like feature map that uses mismatches to capture inexact string matching. The feature space is indexed by all possible k-length subsequences $\alpha = a_1 a_2 \ldots a_k$, where each a_i is a character in the alphabet $\mathcal{A}$ of amino acids. The feature map is defined on k-gram α by $\Phi(\alpha) = (\phi_\beta(\alpha))_{\mathcal{A}^k}$ where $\phi_\beta(\alpha) = 1$ if α is within m mismatches of β, 0 otherwise, and is extended additively to longer sequences: $\Phi(x) = \sum_{k\text{-grams}\alpha \in x} \Phi(\alpha)$. The mismatch kernel can be computed efficiently using a trie data structure: the complexity of calculating $K(x, y)$ is $O(c_K(|x| + |y|))$, where $c_K = k^{m+1}|\mathcal{A}|^m$. For typical kernel parameters $k = 5$ and $m = 1$ (Leslie et al., 2002), the mismatch kernel is fast, scalable and yields impressive performance.

Other direct string kernel methods include pair HMM and convolution kernels (Watkins, 1999; Haussler, 1999; Lodhi et al., 2000), which are quite general but also have complexity $O(|x||y|)$; more recent and related string alignment kernels (Saigo et al., 2004), also with complexity $O(|x||y|)$; and exact-matching string kernels built with suffix trees and suffix links, with complexity $O(|x| + |y|)$ (Vishwanathan and Smola, 2002). Inexact string matching models similar to the mismatch kernel but with complexity $O(c_K(|x| + |y|))$, with c_K independent of alphabet size, have also been presented (Leslie and Kuang, 2003). The motif kernel (Ben-Hur and Brutlag, 2003) uses features that are built from a fixed database of motifs; computing these features is linear in the length of the sequence. Finally, almost all these kernels can be constructed using the rational kernel framework of Cortes et al. (2002). We concentrate on the mismatch kernel representation for the current work.

3 SEMI-SUPERVISED KERNELS FOR PROTEIN SEQUENCES

In semi-supervised learning, one tries to improve a classifier trained on labeled data by exploiting a relatively large set of unlabeled data. An extensive review of techniques can be found in Seeger (2001). It has been shown experimentally that under certain conditions, the decision function

can be estimated more accurately in a semi-supervised setting, yielding lower generalization error. The most common assumption one makes in this setting is called the "cluster assumption," namely that *the class does not change in regions of high density*. Equivalently, one assumes that the true decision boundary lies in regions of low density. In this section, we review previous semi-supervised kernel approaches that implement the cluster assumption.

3.1 Cluster kernels

We will focus on classifiers that re-represent the given data to reflect structure revealed by unlabeled data. The main idea is to change the distance metric so that the relative distance between two points is smaller if the points are in the same cluster. If one is using kernels, rather than explicit feature vectors, one can modify the kernel representation by constructing a cluster kernel.

Previous work of Chapelle et al. (2002) presented a general framework for producing cluster kernels by modifying the eigenspectrum of the kernel matrix. Two of the main methods presented are the *random walk kernel* and the *spectral clustering kernel*. The random walk kernel is a normalized and symmetrized version of a transition matrix corresponding to a t-step random walk. As described in Szummer and Jaakkola (2001), one can define a random walk representation by viewing an RBF kernel as a transition matrix of a random walk on a graph with vertices x_i. One then uses the eigendecomposition of the normalized transition matrix to compute the t-step random walk kernel. The spectral clustering kernel is a simple use of the representation derived from spectral clustering (Ng et al., 2001) using the first k eigenvectors of a normalized affinity matrix.

A serious problem with these methods is that one must diagonalize a matrix of size m, where m is the number of labeled and unlabeled data, giving a complexity $O(m^3)$. Other methods of implementing the cluster assumption such as transductive SVMs (Joachims, 1999) also suffer from computational efficiency issues. A second drawback is that these kernels are better suited to a *transductive* setting (where one is given both the unlabeled and test points in advance) rather than a semi-supervising setting. In order to estimate the kernel for a sequence not present during training, one is forced to solve a difficult regression problem (Chapelle et al., 2002). In Sections 4 and 5, we will describe two simple methods to implement the cluster assumption that do not suffer from these issues.

3.2 Model-based semi-supervised kernels

Another approach for exploiting unlabeled data in kernel methods is to use semi-supervised training of probabilistic models and then define model-based kernels. For example, in the Fisher kernel approach (Jaakkola et al., 2000), one can use an iterative training algorithm to alternately pull in homologs of the positive training examples from a large unlabeled sequence database and reestimate the profile HMM.

More recently, Kuang et al. (2004) introduced a semi-supervised profile-based string kernel for protein sequences. In the profile kernel approach, each sequence is represented by a profile estimated from a large sequence database (for example, using PSI-BLAST), and each length k segment of the profile is used to define the local mutation neighborhood and a corresponding contribution to the feature vector in a k-mer feature space. Unlike the Fisher kernel approach—where a single probabilistic model is used to define feature vectors for sequence examples—in the profile kernel, every example is represented by a probabilistic model in the form of a profile, and the kernel is defined on profile examples. The profile kernel method achieves state-of-the-art performance for the remote homology detection task, strongly outperforming all strictly supervised methods, including the mismatch kernel, the SVM-pairwise method, string alignment kernels (Saigo et al., 2004), and other recent methods (Kuang et al., 2005). We compare the new cluster kernel methods we define here to the profile kernel method in large scale benchmark experiments in Section 6.3.

4 THE NEIGHBORHOOD MISMATCH KERNEL

In this section and the next, we introduce two fast and general cluster kernels that leverage unlabeled data to improve a base kernel representation. Unlike other cluster kernel approaches, our kernels make use of two complementary (dis)similarity measures: (1) a base kernel representation that implicitly exploits features useful for discrimination between classes; and (2) a distance measure that describes how close examples are to each other. In our application to protein classification, we use the mismatch string kernel as the base kernel and standard sequence comparison metrics (such as BLAST or PSI-BLAST E-values) as the distance measure. We note that string kernels have proved to be powerful representations for SVM classification (Leslie et al., 2002) but do not give sensitive pairwise similarity scores like the BLAST family methods; thus the two sequence similarity measures play distinct roles in the kernel definition.

For our first cluster kernel, we use a standard sequence similarity measure like BLAST or PSI-BLAST to define a neighborhood for each input sequence. The neighborhood $\text{Nbd}(x)$ of sequence x is the set of sequences x' with similarity score to x below a fixed E-value threshold, together with x itself. Now given a fixed original feature representation, we represent x by the average of the feature vectors for members of its neighborhood: $\Phi_{nbd}(x) = \frac{1}{|\text{Nbd}(x)|} \sum_{x' \in \text{Nbd}(x)} \Phi_{orig}(x')$. The neighborhood kernel is

then defined by:

$$K_{nbd}(x, y) = \frac{\sum_{x' \in \text{Nbd}(x), y' \in \text{Nbd}(y)} K_{orig}(x', y')}{|\text{Nbd}(x)||\text{Nbd}(y)|}.$$

We will see in the experimental results that this simple neighborhood-averaging technique, used in a semi-supervised setting with the mismatch kernel, dramatically improves classification performance.

In general, computing each neighborhood kernel value is quadratic in neighborhood size, as is clear from the kernel expression given above. However, in the special case where we use the mismatch kernel as base kernel, we can modify the mismatch kernel algorithm by presenting each neighborhood set as a concatenation of the neighbor sequences (keeping track of where the ends of sequences are located); using a trie data structure, the kernel computation is linear in sequence length, giving a complexity of $O(k^{m+1}|\mathcal{A}|^m(\sum x' + \sum y'))$ (that is, linear in neighborhood size) to compute $K_{nbd}(x, y)$, where $|\mathcal{A}|$ is the size of the alphabet of amino acids (Leslie et al., 2002).

To see how the neighborhood approach fits with the cluster assumption, consider a set of points in feature space that form a "cluster" or dense region of the data set, and consider the region R formed by the union of the convex hulls of the neighborhood point sets. If the dissimilarity measure is a true distance, then the neighborhood averaged vector $\Phi_{nbd}(x)$ stays inside the convex hull of the vectors in its neighborhood, and all the neighborhood vectors stay within region R. In general, the cluster contracts inside R under the averaging operation. Thus, under the new representation, different clusters can become better separated from each other.

5 THE BAGGED MISMATCH KERNEL

A number of existing clustering techniques are much more efficient than the methods mentioned in Section 3. For example, the classical k-means algorithm is $O(rkmd)$, where m is the number of data points, d is their dimensionality, and r is the number of iterations required. Empirically, this running time grows sublinearly with k, m and d. Therefore, in practice, it is computationally efficient to run k-means multiple times, which can be useful because k-means can converge to local minima. We therefore consider the following method:

1. Run k-means n times, giving $p = 1, \ldots, n$ cluster assignments $c_p(x_i)$ for each i.

2. Build a bagged-clustering representation based upon the fraction of times that x_i and x_j are in the same cluster:

$$K_{bag}(x_i, x_j) = \frac{\sum_p [c_p(x_i) = c_p(x_j)]}{n}. \qquad (1)$$

3. Take the product between the original and bagged kernel:

$$K(x_i, x_j) = K_{orig}(x_i, x_j) \cdot K_{bag}(x_i, x_j)$$

Because k-means gives different solutions on each run, step (1) will give different results; for other clustering algorithms one could sub-sample the data instead. Step (2) is a valid kernel because it is the inner product in an nk-dimensional space $\Phi(x_i) = \langle [c_p(x_i) = q] : p = 1, \ldots, n, q = 1, \ldots, k \rangle$, and products of kernels as in step (3) are also valid kernels. The intuition behind the approach is that the original kernel is rescaled by the "probability" that two points are in the same cluster, hence encoding the cluster assumption. To estimate the kernel on a test sequence x in a semi-supervised setting, one can assign x to the nearest cluster in each of the bagged runs to compute $K_{bag}(x, x_i)$. We apply the bagged kernel method with K_{orig} as the mismatch kernel and K_{bag} built by running k-means on the distances induced by PSI-BLAST.

6 EXPERIMENTS

We measure the recognition performance of cluster kernel methods by testing their ability to classify protein domains into superfamilies in the Structural Classification of Proteins (SCOP) (Murzin et al., 1995). For the purposes of this experiment, two domains that come from the same superfamily are assumed to be homologous, and two domains from different folds are assumed to be unrelated. For pairs of proteins in the same fold but different superfamilies, their relationship is uncertain, and so these pairs are not used in evaluating the algorithm. This labeling scheme has been used in several previous studies of remote homology detection algorithms (Jaakkola et al., 2000; Liao and Noble, 2002). We use the same 54 target families and the same test and training set splits as in the remote homology experiments in Liao and Noble (2002). The sequences are 7,329 SCOP domains obtained from version 1.59 of the database after purging with astral.stanford.edu so that no pair of sequences share more than 95% identity. Compared to Liao and Noble (2002), we reduce the number of available labeled training patterns by roughly a third. Data set sequences that were neither in the training nor test sets for experiments from Liao and Noble (2002) are included as unlabeled data.

All methods are evaluated using receiver operating characteristic (ROC) analysis (Hanley and McNeil, 1982). An ROC curve plots the rate of true positives as a function of the rate of false positives at varying decision thresholds. The ROC score is the area under this curve. A perfect classifier, which places all positive examples above all negative examples, receives an ROC score of 1, and a random classifier receives a score of approximately 0.5. In addition to the ROC score, we compute the ROC$_{50}$ score, which is the ROC score computed only up to the first 50 false positives (Gribskov and Robinson, 1996). This score focuses on the top of the ranking, which in some applications is the most important.

In all experiments, we use an SVM classifier with a small soft margin parameter, set as $K_{ii} \leftarrow K_{ii} + \gamma$ where γ is 0.02ρ times the median diagonal kernel entry, and ρ is the fraction

of training set sequences that have the same label as the i^{th} sequence. The SVM computations are performed using the freely available Spider Matlab machine learning package.

6.1 Semi-supervised setting

Our first experiment shows that the neighborhood mismatch kernel makes better use of unlabeled data than the baseline method of "pulling in homologs" prior to training the SVM classifier, that is, simply finding close homologs of the positive training examples in the unlabeled set and adding them to the positive training set for the SVM. Homologs come from the unlabeled set (not the test set), and "neighbors" for the neighborhood kernel come from the training plus unlabeled data. We compare the methods using the mismatch kernel representation with $k = 5$ and $m = 1$, as used in Leslie et al. (2002). Homologs are chosen via BLAST or PSI-BLAST as having a pairwise E-value less than 0.05 (the default parameter setting (Altschul et al., 1990)) with any of the positive training samples. The neighborhood mismatch kernel uses the same threshold to choose neighborhoods. For the neighborhood kernel, we normalize before and after the averaging operation via $K_{ij} \leftarrow K_{ij}/\sqrt{K_{ii}K_{jj}}$. The results are given in Figure 1 and Table 1.

	BLAST		PSI-BLAST	
	ROC$_{50}$	ROC	ROC$_{50}$	ROC
mismatch kernel	0.416	0.870	0.416	0.870
mismatch kernel + homologs	0.480	0.900	0.550	0.910
neighborhood mismatch kernel	0.639	0.922	0.699	0.923

Table 1. Mean ROC$_{50}$ and ROC scores over 54 target families for *semi-supervised* experiments, using BLAST and PSI-BLAST for adding homologs and defining the neighborhood kernel.

Figure 1 plots the number of families achieving a given ROC$_{50}$ score. A signed rank test shows that the neighborhood mismatch kernel yields significant improvement over adding homologs (p-value 3.9e-05). Note that the PSI-BLAST scores in these experiments are built using the whole database of 7,329 sequences (that is, test sequences in a given experiment are also available to the PSI-BLAST algorithm), so these results are slightly optimistic. However, the comparison of methods in a truly inductive setting using BLAST shows the same improvement of the neighborhood mismatch kernel over adding homologs (p-value 8.4e-05).

The improvement from the neighborhood kernel does not come from the BLAST and PSI-BLAST representations alone: the mean ROC$_{50}$ score for these representations using an empirical map (see the transductive setting for a description) are 0.368 and 0.533 respectively without pulling in homologs, and 0.448 and 0.595 with pulled in homologs. Moreover, simply adding the BLAST and mismatch kernels together (using an empirical map) without using homologs yields a mean ROC$_{50}$ of 0.3943, so it is also not because

the methods give independent information about the targets which can be easily combined.
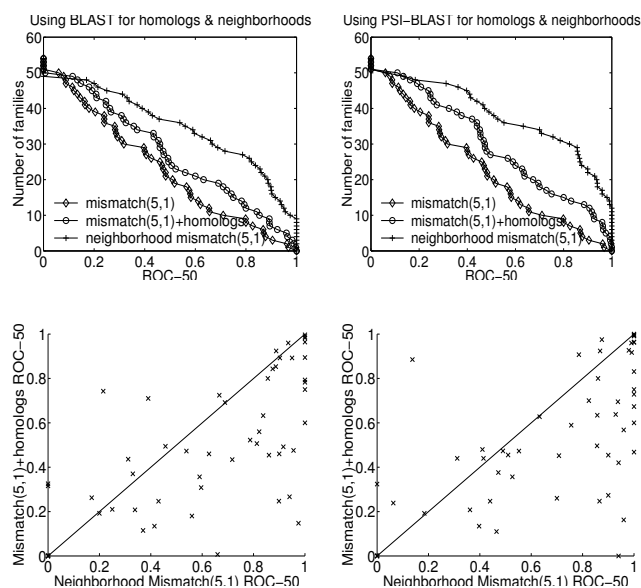

Fig. 1. Comparison of protein representations and classifiers using unlabeled data. The mismatch kernel is used to represent proteins, with close homologs being pulled in from the unlabeled set with BLAST (left) or PSI-BLAST (right). Building a neighborhood with the neighborhood mismatch kernel in both cases improves over the baseline of pulling in homologs. *Note:* We also pull in homologs during the SVM training for the neighborhood kernel.

6.2 Transductive setting

In the following experiments, we consider a *transductive* setting, in which the test points are given to the methods in advance as unlabeled data, giving slightly improved results over the last section. Although this setting is unrealistic for a real protein classification system, it enables comparison with random walk and spectral clustering kernels, which do not easily work in another setting. In Figure 2 (left), we again show the mismatch kernel compared with pulling in homologs and the neighborhood kernel. This time we also compare with the bagged mismatch kernel using bagged k-means with $k = 100$ and $n = 100$ runs, which gave the best results. We found the method quite insensitive to k. The result for $k = 400$ is also given in Table 2. We then compare these methods to using random walk and spectral clustering kernels. Both methods do not work well for the mismatch kernel (see online supplement), perhaps because the feature vectors are so orthogonal. However, for a PSI-BLAST representation via empirical kernel map, the random walk outperforms pulling in homologs. We take the empirical map with $\Phi(x) = (\exp(-\lambda d(x_1, x)), \ldots, \exp(-\lambda(d(x_m, x))))$, where $d(x, y)$ are PSI-BLAST E-values and $\lambda = \frac{1}{1000}$, which

improves over a linear map. We report results for the best parameter choices, $t = 2$ for the random walk and $k = 200$ for spectral clustering. We found the latter quite brittle with respect to the parameter choice; results for other parameters can be found on the supplemental web site. For pulling in close homologs, we take the empirical kernel map only for points in the training set and the chosen close homologs. Finally, we also run transductive SVMs. The results are given in Table 2 and Figure 2 (right). A signed rank test (with adjusted p-value cut-off of 0.05) finds no significant difference between the neighborhood kernel, the bagged kernel ($k = 100$), and the random walk kernel in this transductive setting. Thus the new techniques are comparable with random walk, but are feasible to calculate on full scale problems.


Fig. 2. Comparison of protein representations and classifiers using unlabeled data in a *transductive* setting. Neighborhood and bagged mismatch kernels outperform pulling in close homologs (top) and equal or outperform previous semi-supervised methods (bottom). *Note:* We also pull in homologs during the SVM training for the neighborhood and bagged kernels.

	ROC$_{50}$	ROC
mismatch kernel	0.416	0.875
mismatch kernel + homologs	0.625	0.924
neighborhood mismatch kernel	0.704	0.917
bagged mismatch kernel ($k = 100$)	0.719	0.943
bagged mismatch kernel ($k = 400$)	0.671	0.935
PSI-BLAST kernel	0.533	0.866
PSI-BLAST+homologs kernel	0.585	0.873
spectral clustering kernel	0.581	0.861
random walk kernel	0.691	0.915
transductive SVM	0.637	0.874

Table 2. Mean ROC$_{50}$ and ROC scores over 54 target families for *transductive* experiments.

6.3 Large scale experiments

Semi-supervised and transductive methods are most interesting and potentially give greatest benefit in the realistic setting where a large amount of unlabeled data is used. We therefore test our cluster kernel methods in large scale experiments, using 101,602 Swiss-Prot protein sequences as additional unlabeled data. For simplicity, we first give results for both the neighborhood and bagged kernels in the transductive setting, that is, in the case where test sequences are available as additional unlabeled examples in all the experiments. Then, for a clean comparison against the profile kernel, we test the neighborhood kernel and the profile kernel in a semi-supervised setting, where the Swiss-Prot database alone is used as the source of unlabeled data.

For the large scale neighborhood mismatch kernel experiments, we first compute the entire SCOP plus Swiss-Prot kernel (108931-by-108931) matrix with mismatch kernel parameters $k = 5$ and $m = 1$. We then apply the neighborhood averaging operation to produce the 7329-by-7329 kernel matrix for SCOP sequences needed for SVM training. We normalize the kernel matrix before and after the neighborhood averaging operation. Results in Table 3 clearly show that the inclusion of a large amount of additional unlabeled data from Swiss-Prot significantly improves classification performance. Moreover, the neighborhood kernel again outperforms the baseline method of adding homologs of the positive training sequences to the training set.

For the large scale bagged mismatch kernel experiments, the fact that many of the sequences in the Swiss-Prot database are multi-domain protein sequences complicates the clustering step: since the PSI-BLAST E-values used as the dissimilarity metric are based on local alignment, a multi-domain sequence can be similar to many unrelated single-domain sequences, and hence the clustering algorithm may fail to converge. As an approximate remedy, we only use Swiss-Prot protein sequences with maximal length of 250 for the large scale k-means clustering, reasoning that most multi-domain sequences would be eliminated by this length constraint. We randomly sample 30,000 protein sequences from the set of Swiss-Prot with length 250 or less to use as unlabeled data for clustering. Since the method mainly depends on the quality of the clusters containing the labeled points, we terminate the k-means clustering algorithm once there are no more changes in the label assignment for the SCOP sequences. It is worth noting that a small amount of 2-domain sequences may have length below our cut-off, but we observe that the k-means clustering algorithm still behaves relatively stably. We use the same mismatch kernel parameters for the bagged kernel as the ones we use for the small scale bagged kernel experiments. A comparison of results is shown in Table 3. Again, bagged kernel performance significantly improves when a large amount of unlabeled data is provided to the clustering algorithm.

Finally, we compare the performance of the stronger of the cluster kernels, the neighborhood kernel, to a state-of-the-art semi-supervised kernel method, the profile kernel. The profile kernel representation depends on estimating sequence profiles for each input sequence using a large sequence database, and therefore we only present results in the large-scale

	Without Swiss-Prot		With Swiss-Prot	
	ROC_{50}	ROC	ROC_{50}	ROC
mismatch kernel + homologs	0.625	0.924	0.706	0.945
neighborhood mismatch kernel	0.704	0.917	0.871	0.971
bagged mismatch kernel ($k = 100$)	0.719	0.943	0.803	0.953
bagged mismatch kernel ($k = 400$)	0.671	0.935	0.775	0.955

Table 3. Mean ROC_{50} and ROC scores over 54 target families for *large scale transductive* experiments. *Note:* We include homologs from the unlabeled set and the test set (SCOP+Swiss-Prot) for the training of all our SVMs.

setting. In these experiments, we use a semi-supervised training set-up: the Swiss-Prot database alone is used as the source of unlabeled data for estimating PSI-BLAST profiles and defining sequence neighborhoods; SCOP sequences are not used for profile learning or for neighborhood averaging. For the cleanest comparison, we do not add SCOP homologs to the positive training set before training the SVMs. Mean ROC_{50} and ROC results are given in Table 4, and a comparison of ROC_{50} results over all experiments is given in Figure 3. While the cluster kernel method does not outperform the profile kernel on average, results of the two methods are similar (20 wins, 25 losses, 9 ties for the cluster kernel); a signed rank test with a p-value threshold of 0.05 finds no significant difference in performance between the two methods.

The computational cost of the profile kernel depends on the parameter σ controlling the size of the local mutation neighborhoods; the mutation neighborhood of each length k window of the profile consists of k-mers whose negative log likelihood given the profile is less than σ. If M_σ is the maximum size of any local mutation neighborhood in the input sequences, the complexity of computing the profile kernel for profiles of sequences x and y can be bounded by $O(kM_\sigma(|x| + |y|))$. In practice, in the large-scale experiments reported, the profile kernel takes less time to compute than the neighborhood kernel, since we do not limit the number of sequences represented in the neighborhoods and hence the neighborhood representation is much less compact than the profile kernel representation. Various schemes could make the neighborhood kernel running time comparable or faster than the profile kernel, for example, sampling from the neighborhood rather than using all sequences; one would have to investigate how much sampling is needed to retain classification performance.

While the neighborhood and profile kernels have a similar number of wins in this set of experiments, performance on individual experiments can be quite different. We speculate that the profile kernel might perform best when the training profiles themselves are not too distant from the test sequences (as measured, for example, by PSI-BLAST E-value), while the neighborhood kernel might perform better in the more difficult situation where the test sequences are poorly described by every positive training profile. (Note that, in both situations, the discriminative profile kernel SVM

outperforms the PSI-BLAST algorithm used directly as a ranking method.) To test this hypothesis, we consider the top 10 wins of the profile kernel method and of the neighborhood kernel method, as ranked by the difference in ROC_{50} performance, and for each positive test sequence in these experiments, we find the positive training profile that detected the test sequence with minimum PSI-BLAST E-value. We then rank the test sequences by this minimum E-value and record the median E-value in this list. We find that in the top 10 profile kernel wins, 8 of the experiments has median minimum E-value less than .05, and 9 are less than .1, while in the neighborhood kernel wins, only 2 experiments have median minimum E-value less than .05, and only 3 are less that .1. We conclude that the neighborhood kernel may have an advantage over the profile kernel in situations where the training profiles are very distant from the remote homolog sequences to be detected.

	ROC_{50}	ROC
neighborhood mismatch kernel	0.810	0.955
profile kernel	0.842	0.980

Table 4. Mean ROC_{50} and ROC scores over 54 target families for *large scale semi-supervised* experiments. *Note:* We do not include homologs from the unlabeled set (Swiss-Prot) for the training of our SVMs in these experiments.
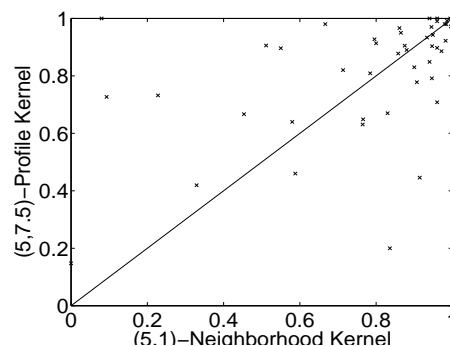

Fig. 3. Comparison of neighborhood kernel and profile kernel ROC_{50} performance for *large scale semi-supervised* experiments. No homologs were added to the training set for the purpose of training the SVMs.

7 DISCUSSION

Two of the most important issues in protein classication are representation of sequences and handling unlabeled data. Two developments in recent kernel methods research, string kernels and cluster kernels, address these issues separately. We have described two kernels — the *neighborhood mismatch kernel* and the *bagged mismatch kernel* — that

combine both approaches and yield state-of-the-art performance in protein classification. These approaches, used with an efficient string kernel, are fast and scalable cluster kernels for sequence data and do not require diagonalization of the kernel matrix as in other cluster kernel methods. A potential direction for improvement in the neighborhood kernel would be to extract only those segments of "neighboring" sequences that correspond to the local alignment-based E-value score; when we use entire multi-domain Swiss-Prot sequences as neighbors of a single-domain SCOP sequence, these neighbor sequences may include long regions that are unrelated to the SCOP domain, and hence we introduce noise in the neighborhood averaging operation.

While we have motivated our kernels by earlier work on cluster kernels and the cluster assumption, one can also view the neighborhood and bagged kernels as using unlabeled data locally (from nearby sequences or the local cluster) for smoothing the kernel representation. Related work using probabistic models instead of unlabeled data for smoothing includes the recently introduced Bhattacharyya kernel (Jebara et al., 2004), which assigns a probability distribution to each example and defines a kernel on these distributions.

We also compared to the profile-based string kernels of Kuang et al. (2004), which are also based on a semi-supervised learning paradigm. These string kernels are also scalable and achieve very high classification accuracy; in our experiments, the neighborhood kernel performs similarly to the profile kernel. However, the profile kernel method requires producing a profile for each query sequence, which is necessarily tied to alignment. In contrast, the cluster kernels that we present here are more general, in that any dissimilarity measure can be used for neighborhood averaging or bagging and any base kernel chosen for the initial representation. These kernels may therefore be applicable to a wider range of problems. For example, one could use expression coherence in a set of microarray experiments as a measure of functional similarity of genes combined with a base kernel to define cluster kernels for functional gene classification. One could also hope to further improve performance for the protein classification task by using a more powerful base kernel than the mismatch kernel (for example, the string alignment kernel of Saigo et al. (2004)), though the computational expense of the improved base kernel representation may become a concern.

ACKNOWLEDGEMENT

REFERENCES

Altschul, S. F., W. Gish, W. Miller, E. W. Myers, and D. J. Lipman (1990). A basic local alignment search tool. *Journal of Molecular Biology 215*(3), 403–410.

Altschul, S. F., T. L. Madden, A. A. Schaffer, J. Zhang, Z. Zhang, W. Miller, and D. J. Lipman (1997). Gapped BLAST and PSI-BLAST: A new generation of protein database search programs. *Nucleic Acids Research 25*, 3389–3402.

Ben-Hur, A. and D. Brutlag (2003). Remote homology detection: a motif based approach. *Bioinformatics 19*, i26–i33.

Chapelle, O., J. Weston, and B. Schölkopf (2002). Cluster kernels for semi-supervised learning. *Advances in Neural Information Processing Systems 15*, 601–608.

Cortes, C., P. Haffner, and M. Mohri (2002). Rational kernels. *Advances in Neural Information Processing Systems 15*.

Gribskov, M. and N. L. Robinson (1996). Use of receiver operating characteristic (roc) analysis to evaluate sequence matching. *Computers and Chemistry 20*(1), 25–33.

Hanley, J. A. and B. J. McNeil (1982). The meaning and use of the area under a receiver operating characteristic (ROC) curve. *Radiology 143*, 29–36.

Haussler, D. (1999). Convolution kernels on discrete structure. Technical report, University of California, Santa Cruz.

Jaakkola, T., M. Diekhans, and D. Haussler (2000). A discriminative framework for detecting remote protein homologies. *Journal of Computational Biology 7*(1–2), 95–114.

Jebara, T., R. Kondor, and A. Howard (2004). Probability product kernels. *Journal of Machine Learning 5*, 819–844.

Joachims, T. (1999). Transductive inference for text classification using support vector machines. *Proceedings of Intenational Conference on Machine Learning*, 200–209.

Krogh, A., M. Brown, I. Mian, K. Sjolander, and D. Haussler (1994). Hidden markov models in computational biology: Applications to protein modeling. *Journal of Molecular Biology 235*, 1501–1531.

Kuang, R., E. Ie, K. Wang, K. Wang, M. Siddiqi, Y. Freund, and C. Leslie (2004). Profile-based string kernels for remote homology detection and motif extraction. In *3rd International IEEE Computer Society Computational Systems Bioinformatics Conference*, pp. 152–160. IEEE Computer Society.

Kuang, R., E. Ie, K. Wang, K. Wang, M. Siddiqi, Y. Freund, and C. Leslie (2005). Profile kernels for detecting remote protein homologs and discriminative motifs. *Journal of Bioinformatics and Computational Biology*. To appear.

Leslie, C., E. Eskin, J. Weston, and W. S. Noble (2002). Mismatch string kernels for SVM protein classification. *Advances in Neural Information Processing Systems 15*, 1441–1448.

Leslie, C. and R. Kuang (2003). Fast kernels for inexact string matching. *Sixteenth Annual Conference on Learning Theory and Seventh Kernel Workshop*, 114–128.

Liao, C. and W. S. Noble (2002). Combining pairwise sequence similarity and support vector machines for remote protein homology detection. *Proceedings of the Sixth Annual International Conference on Research in Computational Molecular Biology*, 225–232.

Lodhi, H., J. Shawe-Taylor, N. Cristianini, and C. Watkins (2000). Text classification using string kernels. In *Advances in Neural Information Processing Systems 13*, pp. 563–569.

Murzin, A. G., S. E. Brenner, T. Hubbard, and C. Chothia (1995). SCOP: A structural classification of proteins database for the investigation of sequences and structures. *Journal of Molecular Biology 247*(4), 536–540.

Ng, A., M. Jordan, and Y. Weiss (2001). On spectral clustering: analysis and an algorithm. *Advances in Neural Processing Information Systems 14*, 849–856.

Park, J., K. Karplus, C. Barrett, R. Hughey, D. Haussler, T. Hubbard, and C. Chothia (1998). Sequence comparisons using multiple sequences detect twice as many remote homologues as pairwise methods. *Journal of Molecular Biology 284*(4), 1201–1210.

Saigo, H., J. Vert, N. Uea, and T. Akutsu (2004). Protein homology detection using string alignment kernels. *Bioinformatics*.

Seeger, M. (2001). Learning with labeled and unlabeled data. Technical report, University of Edinburgh.

Smith, T. and M. Waterman (1981). Identification of common molecular subsequences. *Journal of Molecular Biology 147*(1), 195–197.

Szummer, M. and T. Jaakkola (2001). Partially labeled classification with Markov random walks. *Advances in Neural Information Processing Systems 14*.

Vishwanathan, S. V. N. and A. Smola (2002). Fast kernels for string and tree matching. *Advances in Neural Information Processing Systems 15*, 585–592.

Watkins, C. (1999). Dynamic alignment kernels. Technical report, Royal Holloway, University of London.

Weston, J., C. Leslie, D. Zhou, A. Elisseeff, and W. S. Noble (2003). Cluster kernels for semi-supervised protein classification. *Advances in Neural Information Processing Systems 17*.

Zhu, X. and Z. Ghahramani (2002). Learning from labeled and unlabeled data with label propagation. Technical report, Carnegie Mellon University.